Supplementary Artifact Release Statement

Prepared 12 October 2026. This statement accompanies the code-only Poison Detection artifact. Sections S1-S5 provide the release evidence for the responsible-research checklist. It describes the prepared distribution and does not assert that an upload to a hosting service has already occurred.

S1. Licenses and release scope

Project-authored code and documentation in the prepared repository are distributed under the MIT license, with copyright attributed to the Poison Detection contributors. The full license and third-party notices are included. External software is installed separately, not vendored. Kronfluence is Apache-2.0; its source and license are linked in the notices. Local stability wrappers are identified as project modifications, not an upstream release.

The distribution excludes all original and rewritten corpus records, prompt and completion caches, model/tokenizer weights, learned adapters, per-example scores and predictions, sample-ID lists, logs, external repositories and unsanitized research history. An explicit file inventory governs the repository. A verifier checks its index and every locally reachable commit for excluded paths and the translated baseline; source ZIP exports additionally include SHA-256 hashes.

IMDB's source card leaves detailed licensing information incomplete. The precise task/revision and underlying instance rights of the historical SuperNaturalInstructions-based pool were not fully reconstructed. We therefore do not redistribute those records or their derivatives. SST-2-derived Hidden Killer data is also excluded: the containing repository's current MIT indication does not substitute for the original corpus terms. GSM8K has an upstream MIT license; its records are excluded consistently with the code-only scope.

The primary FLAN-T5-small checkpoint and the referenced T5 LM-adapt and Qwen2.5-7B models identify Apache-2.0 terms. DeepSeek-Coder model weights have separate model terms, and optional Llama examples have their respective community-license conditions. None of these weights is redistributed. The notices give source links and the recorded primary model/dataset revisions. Configurations naming a model do not grant a new license for that model.

No explicit license was found in the upstream released unigram Z-Score repository used by the historical translation. That implementation is archived separately and excluded from this distribution; the MIT notice does not cover it. The public adapter raises an informative error if it is requested. An explicit PD run that omits this baseline uses a separate run tag and does not claim to reproduce the excluded comparison or removal arm. Historical measurements are retained as aggregate records without changing their values.

This scope updates the public-release descriptions in the manuscript's Appendix A and Appendix F: row IDs and per-example records are retained in the local research archive rather than included in the code-only distribution. It does not change the empirical results or claim that any older hosted artifact has been updated.

Sources: Kronfluence license, IMDB source card, Hidden Killer, SuperNaturalInstructions, GSM8K license, released Z-Score. The accompanying THIRD_PARTY_NOTICES.md contains the complete inventory and links.

S2. Personal information and mitigation

Source movie reviews and generated variants may contain identifying information even when author metadata is absent. Public availability does not prove that the text is anonymous, and replacing an example ID does not prevent searching for unchanged text. We do not claim to have fully anonymized the local research data.

An email-format screening pass over two specified local clean training pools identified candidate matches in seven of 10,000 IMDB rows and zero of 6,000 SST-2 rows. These are unadjudicated pattern matches, not confirmed counts of identifiable people. The check does not cover all names, addresses, obfuscations or other identifiers and is not a full-corpus audit. The script, tested file descriptions and aggregate results are provided without matched text or row IDs.

The implemented release mitigation is exclusion of the entire raw/derived text collections, cached prompts, per-example outputs, weights and unsanitized Git history. Only code, documentation and aggregate results are included. Repository checks scan tracked text and historical file versions for credential and email patterns. Local experimental inputs and results are preserved, so the publication preparation does not silently alter the data underlying reported measurements.

S3. Objectionable content and research use

Movie reviews can contain profanity and harmful descriptions, and attack rewriting can preserve or introduce problematic language. A case-insensitive keyword check flagged 49 of the 10,000 screened IMDB rows and zero of the 6,000 screened SST-2 rows. These are keyword candidates, not contextual toxicity judgments; zero matches does not establish absence of objectionable content.

The release excludes the source reviews and all generated/rewritten examples. Documentation warns that local runs can regenerate problematic content. The screening script includes its detection vocabulary, but no source review is published as a worked example. No blanket claim of content filtering or benign generated outputs is made.

The intended use is controlled research on poisoning detection and defenses. Attack-generation code can also facilitate misuse, and a high suspicion score must not be treated as proof of malicious intent. Removing samples can discard useful or minority-distribution data. The work does not establish adaptive robustness or general safety. These use statements do not impose additional restrictions on the MIT software license.

S4. Artifact documentation and coverage

The primary study covers English movie-review sentiment classification on IMDB and SST-2-derived Hidden Killer data, with negative/positive classes and a positive attack target. Historical supplementary work adds English instruction-formatted sentiment and English GSM8K grade-school mathematical word problems. The multilingual ability of a base model is not evaluated-language coverage for this study.

Linguistic interventions include token prefixes, entity-trigger phrases, lexical and negation edits, fixed syntax/style wrappers, released SCPN syntax rewrites, and model-generated style/paraphrase variants. Inspired local implementations are distinguished from author-released generator outputs. Human validation of semantic preservation is not claimed for every rewrite.

No demographic annotations are added for review writers. Their age, gender, ethnicity, nationality and other demographic distributions are unknown, and no representativeness claim is made. Domain, source sampling and sentiment-label selection limit generalization.

The README, data documentation and experiment protocol describe dependencies, source versions, expected input fields, score directions, calibration, seeds, training and evaluation. A synthetic scorer check runs without loading models or datasets. Source acquisition and the missing licensed baseline are explicit prerequisites/limitations. The code-only package is not advertised as a complete reconstruction of undocumented historical mixed-source pools.

S5. Dataset statistics and split construction

The accompanying data_statistics.csv reports exact counts for all 42 main-study runs, including seed, training mode, poison count and evaluation denominators. The following table summarizes the primary pools actually used.

Dataset / pool	Train	Calibration	Clean test	Triggered test	Poison rates
IMDB, small	1,000	1,000	1,000	499	0%, 1%, 5%
IMDB, large	10,000	1,000	1,000	499	0%, 0.5%, 1%, 5%
SST-2 / Hidden Killer	6,000	872	1,819-1,820	910	0%, 1%, 5%

Seeds are 42, 43 and 44. Full fine-tuning covers these rates; LoRA covers the larger pools at 0% and 5%. Pools are balanced before poisoning. Negative-source poisons are relabeled positive, yielding 50, 500 or 300 poisoned rows at 5% for the three sizes. Normalized duplicates and cross-role overlaps are removed. IMDB calibration/test are disjoint subsets of its official test split. SST-2's released dev split supplies calibration, with no additional tuning split. Triggered examples reuse eligible negative-source test examples and are not another independent split. The separate 1,000-example SCPN pilot is excluded from the 42-run count.

For the historical local sentiment files, train/test sizes are 1,000/200, and the training file contains 50 pre-poisoned rows. Observed train labels are 523 positive and 477 negative; test labels are 109 positive and 91 negative. A separate diagnostic file contains 100 original records and three transformation families with two variants each. Historical scripts further subsample these files, so file sizes must not be substituted for every reported experiment's evaluation denominator.

The historical GSM8K removal record uses 7,473 training examples, 374 poisons at 5%, three epochs, and a removal budget of 100. The cached source test file has 1,319 rows; this is not asserted to be the scored test subset. Padding and removal-budget confounds remain. Some supplementary source mappings, splits, seeds, generator settings and thresholds are unavailable, as acknowledged in the original Appendix F and the new data documentation. These gaps are disclosed rather than replaced with invented statistics or pooled primary-study counts.